



Miniboxing

Runtime Specialization on the JVM


Preliminary benchmark results
LAMP, EPFL, 7th of March 2013



Vlad Ureche
vlad.ureche@epfl.ch

Miniboxing transformation

```
class C[@minispec T](implicit mf: Manifest[T]) {  
  var a = new Array[T](5)  
}
```



```
abstract class C[T](implicit mf: Manifest[T]) {  
  def a: Array[T]  
  def a_=(_a: Array[T])  
  def a_!: Array[T]  
  def a_=_(_a: Array[T])  
}
```

Miniboxing transformation

```
class C_J[T] extends C[t]{  
  private[this] _a: Array[T] // erased to Object  
  
  def <init>(T$Type: Byte) = {           // call to the  
    this._a = mbarray_new(T$Type, 5) // library but  
  }                                     // inlined by  
                                     // the backend  
  
  // actual implementations  
  def a: Array[T] = a  
  def a_=(_a: Array[T]) = a = 1( a )  
  def a_J: Array[T] = a_J  
  def a_=_J(_a: Array[T]) = a_=_J( a )  
}
```

T\$type match {
 case INT => new Array[Int](5)
 ...
}.asInstanceOf[Array[T]]

Main benchmark

- `ArrayBuffer.reverse()`
 - VM can still perform array ops
 - **toughest benchmark**
 - **12x speedup over generic**

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

Library Support

- `ArrayBuffer.reverse()`

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

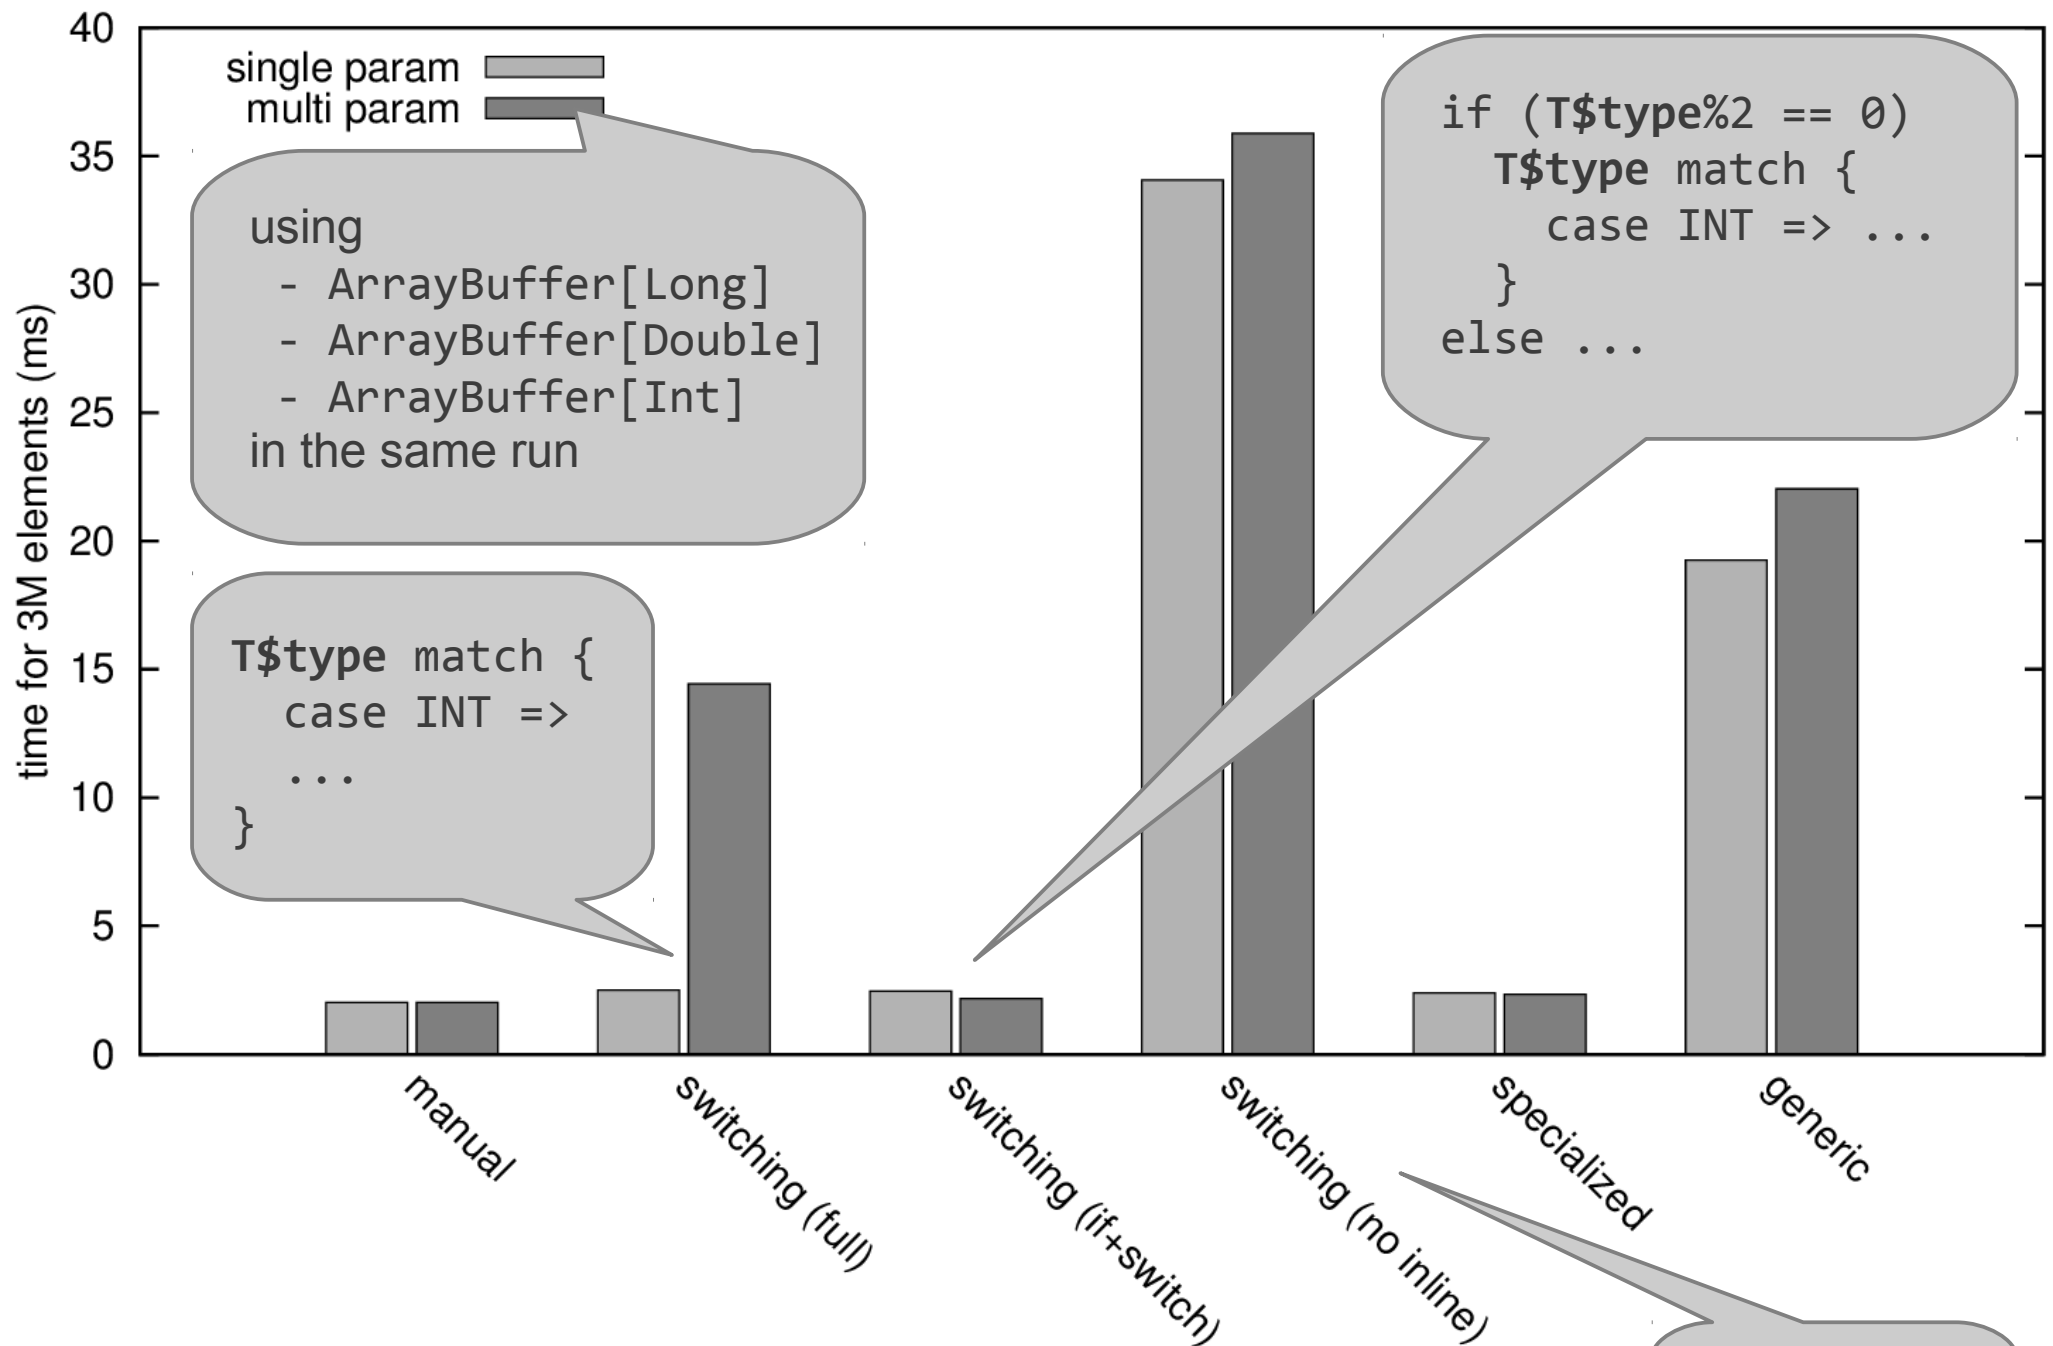
`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

Can the JVM optimize such code?



ArrayBuffer.reverse()

no inlining in
the backend

Library Support

- switch+if works without slowdowns for up to 6 type parameters:

```
if (T$type%2 == 0)
  T$type match {
    case INT => ...    // this path is being benchmarked, fast path
    case CHAR => ...
    case BOOLEAN => ...
    case BYTE => ...
  }
else
  T$type match {
    case DOUBLE => ... // this path has been exercised before, slow path
    case FLOAT => ...  // this path has been exercised before, slow path
    case LONG => ...   // this path has been exercised before, slow path
    case UNIT => ...   // this path has been exercised before, slow path
    case SHORT => ...  // this path has been exercised before, slow path
  }
```

But it's still not a complete solution

Library Support

- Insight:
 - Let's lift the switch on the entire body of reverse
 - The JVM can surely optimize that
 - What if the loop is not in the current method?
 - We'd still be doing the switching at each loop
 - We need a more generic way, instance-wide

Dispatching

- `ArrayBuffer.reverse()`

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

Need to lift the switches outside the loop

Dispatching

- `ArrayBuffer.reverse()`

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val other = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(other)  
    array(index) = tmp2  
    array(other) = tmp1  
    index += 1  
  }  
}
```


`T$Dispatcher.array_get`

`T$Dispatcher.array_get`

`T$Dispatcher.array_update`


`T$Dispatcher.array_update`

We committed to the data representation early

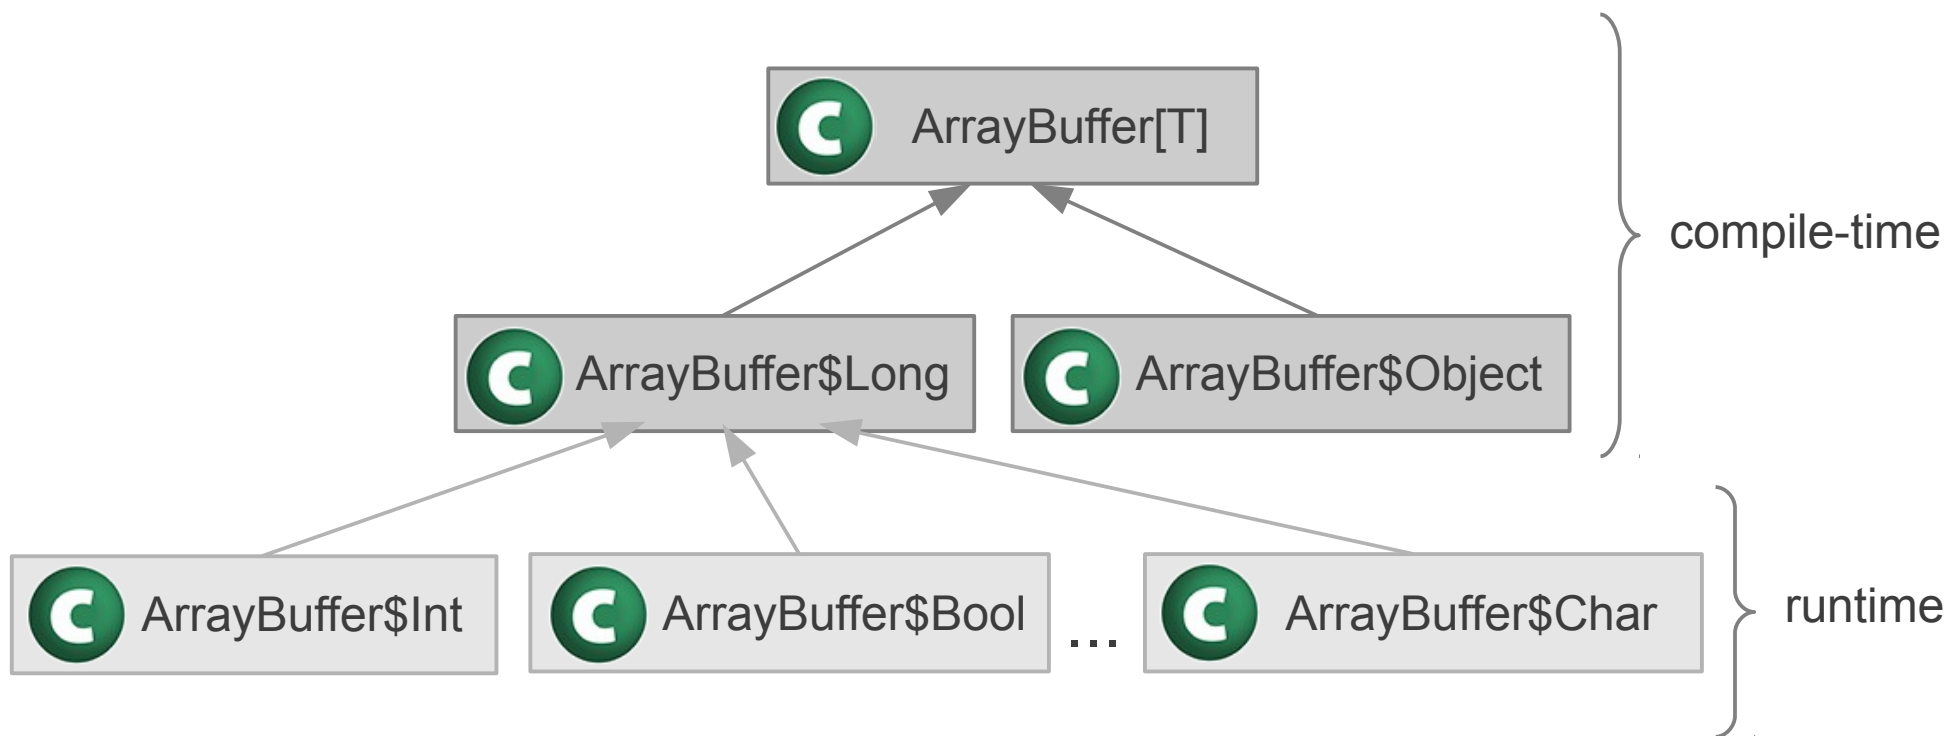


We need to make a promise and stick to it

Runtime Specialization



- dot net-like runtime specialization
 - two stages: compile-time and runtime



What do we gain?

Runtime Specialization


- `ArrayBuffer$Int`
 - has `IntDispatcher` set statically
- Operations can be inlined
 - since `IntDispatcher` is final
 - calls never become megamorphic
 - optimizations can be done inside the loop
- Separate call sites for `array_get`

Class modification and loading cost?


Runtime Specialization

- One-time non-negligible cost
 - load . . . \$Long bytecode
 - modify it
 - build a `Class[_]` object
 - instantiate it
- Apply the cost to the factory
 - instead of the individual class
 - turns out the cost amortizes well (10-50 uses)
 - Only local classloader required

Class modification and loading cost?



`ArrayBuffer.reverse()`




Looks good on this benchmark. What about others?

`ArrayBuffer.reverse()`

All benchmarks

- `ArrayBuffer.insert()`
 - nested calls performance
 - confusing the JVM inline heuristics => this blows up

```
def add(elem: T) = {  
  extend()  
  array(elemCount) = elem  
  elemCount += 1  
}  
  
def extend(): Unit = {  
  if (elemCount == size) {  
    ...  
  }  
}
```




`ArrayBuffer.insert()`

All benchmarks

- `ArrayBuffer.reverse()`
 - VM can still perform array ops
 - **toughest benchmark**
 - **10x+ speedup over generic**

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```




ArrayBuffer.reverse()

All benchmarks

- `ArrayBuffer.contains()`
- `==` operator performance

```
def contains(elem: T): Boolean = {  
    var pos = 0  
    while (pos < elemCount){  
        if (getElement(pos) == elem)  
            return true  
        pos += 1  
    }  
    return false  
}
```

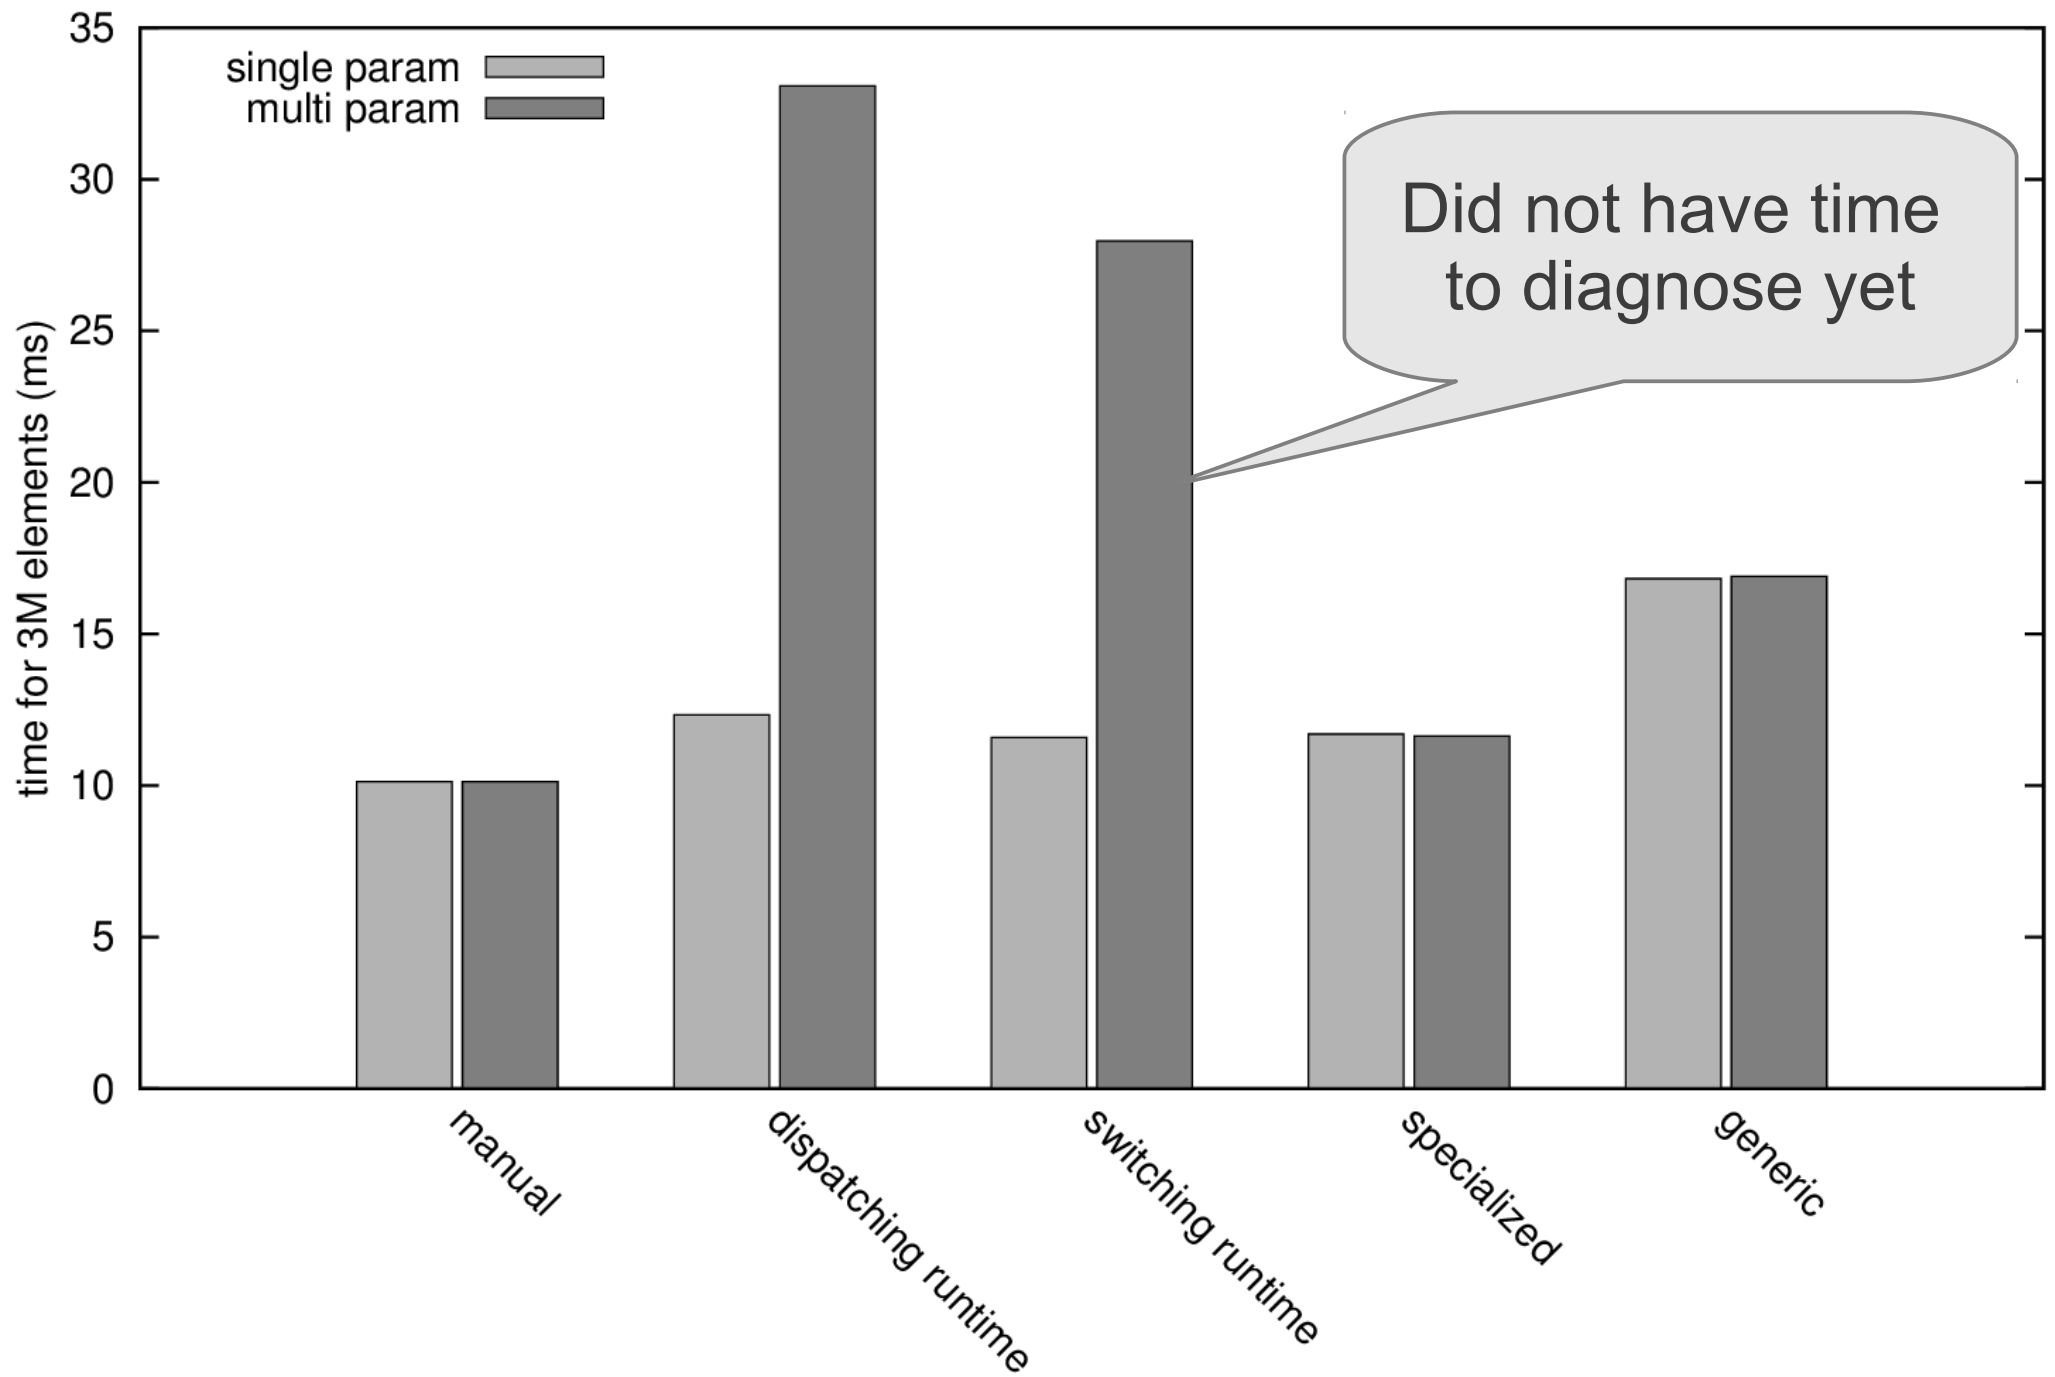


ArrayBuffer.contains()

All benchmarks

- new LinkedList(...)
 - new operator performance
- null tail ends a list

```
class LinkedList[T](val head: T, val tail: LinkedList[T])
```




`new LinkedList(...)`

All benchmarks

- `LinkedList.hashCode()`
 - hashCode performance
 - traversal performance

```
def contains(e: T): Boolean = {  
  @annotation.tailrec def containsTail(list: List[T]):  
Boolean =  
    if (list.head == e)  
      true  
    else if (list.tail == null)  
      false  
    else  
      containsTail(list.tail)  
  
  containsTail(this)  
}
```




LinkedList.hashCode()

All benchmarks

- `LinkedList.contains()`
 - `==` operator performance
 - traversing performance

```
def contains(e: T): Boolean = {  
  @annotation.tailrec def containsTail(list: List[T]):  
Boolean =  
    if (list.head == e)  
      true  
    else if (list.tail == null)  
      false  
    else  
      containsTail(list.tail)  
  
  containsTail(this)  
}
```



LinkedList.contains()

Contributions

- Miniboxing itself, in an open-world assumption
- Exploration of the implementation space
 - Dispatchers
 - Runtime specialization
- Solution that works in practice



Thank you!

github.com/miniboxing



Vlad Ureche
vlad.ureche@epfl.ch