



Miniboxing

Runtime Specialization on the JVM


Under the hood presentation
LAMP, EPFL, 6th of March 2013



Vlad Ureche
vlad.ureche@epfl.ch

Generic Code

- enables reuse and safety
- value types - pointers (Object) vs stack values
- common representation - boxing? NO



We are looking for specialization

Outline

- Generic code
 - Problem
 - State of the art
- Miniboxing
- Evolving miniboxing
- Evaluation

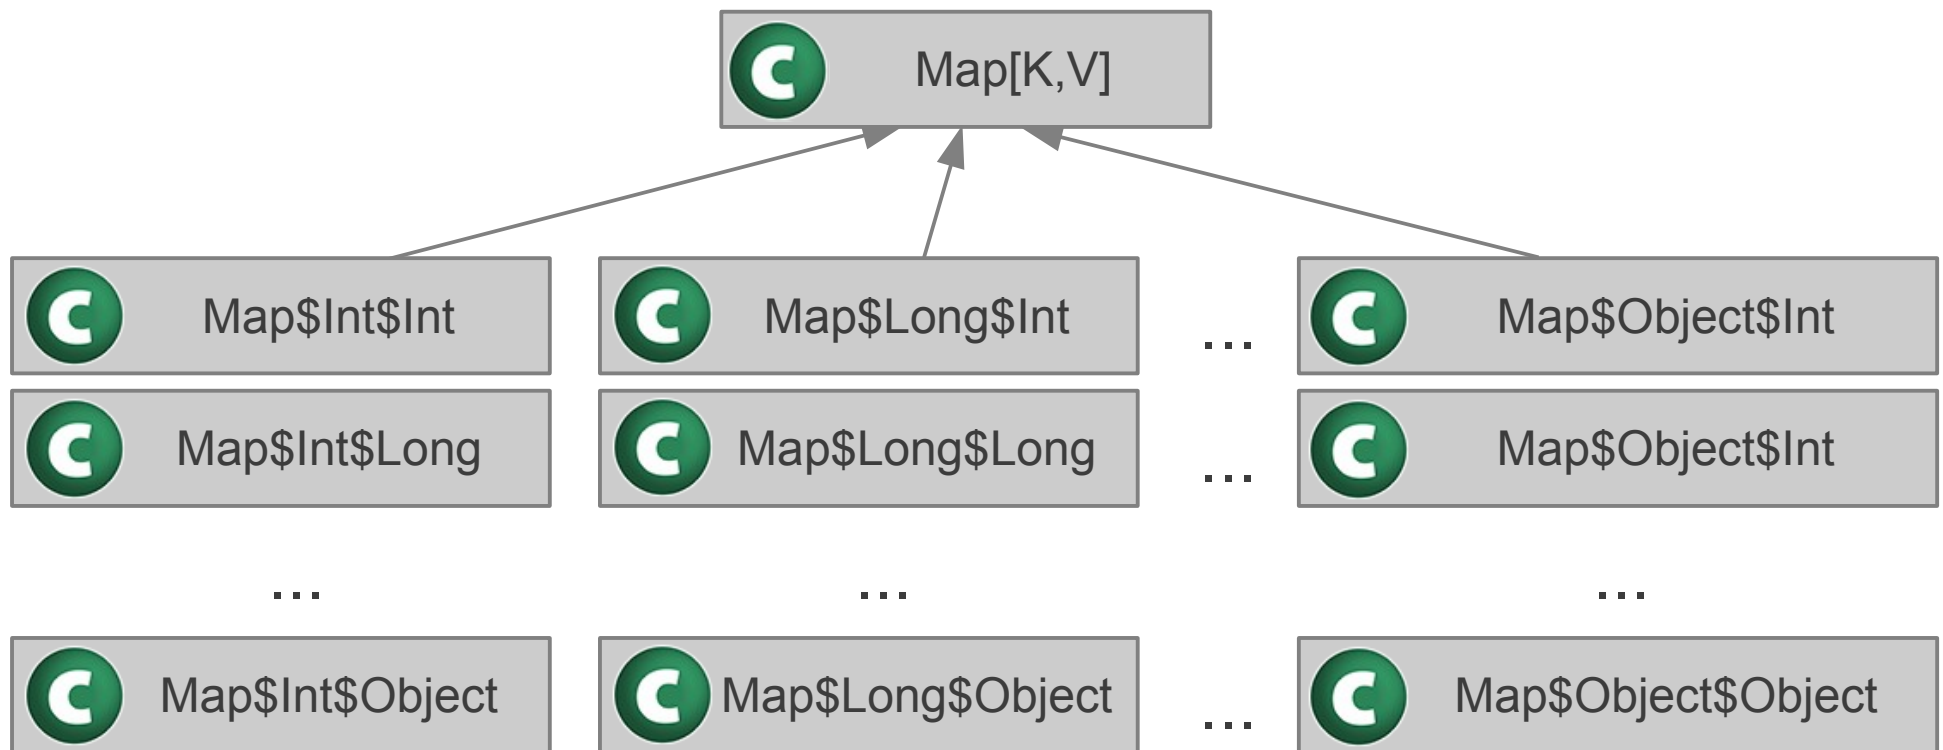
State of the Art

- dot net
 - runtime specialization, in the virtual machine
 - doesn't apply to Java bytecode - erasure
- c++ templates
 - expand at compile time, as needed
 - duplicate instantiations open the door to incompatibilities
- pizza
 - runtime specialization via classloader
 - complex changes, reflection, slow

Hard problem

State of the art - Specialization

- static expansion, in the library

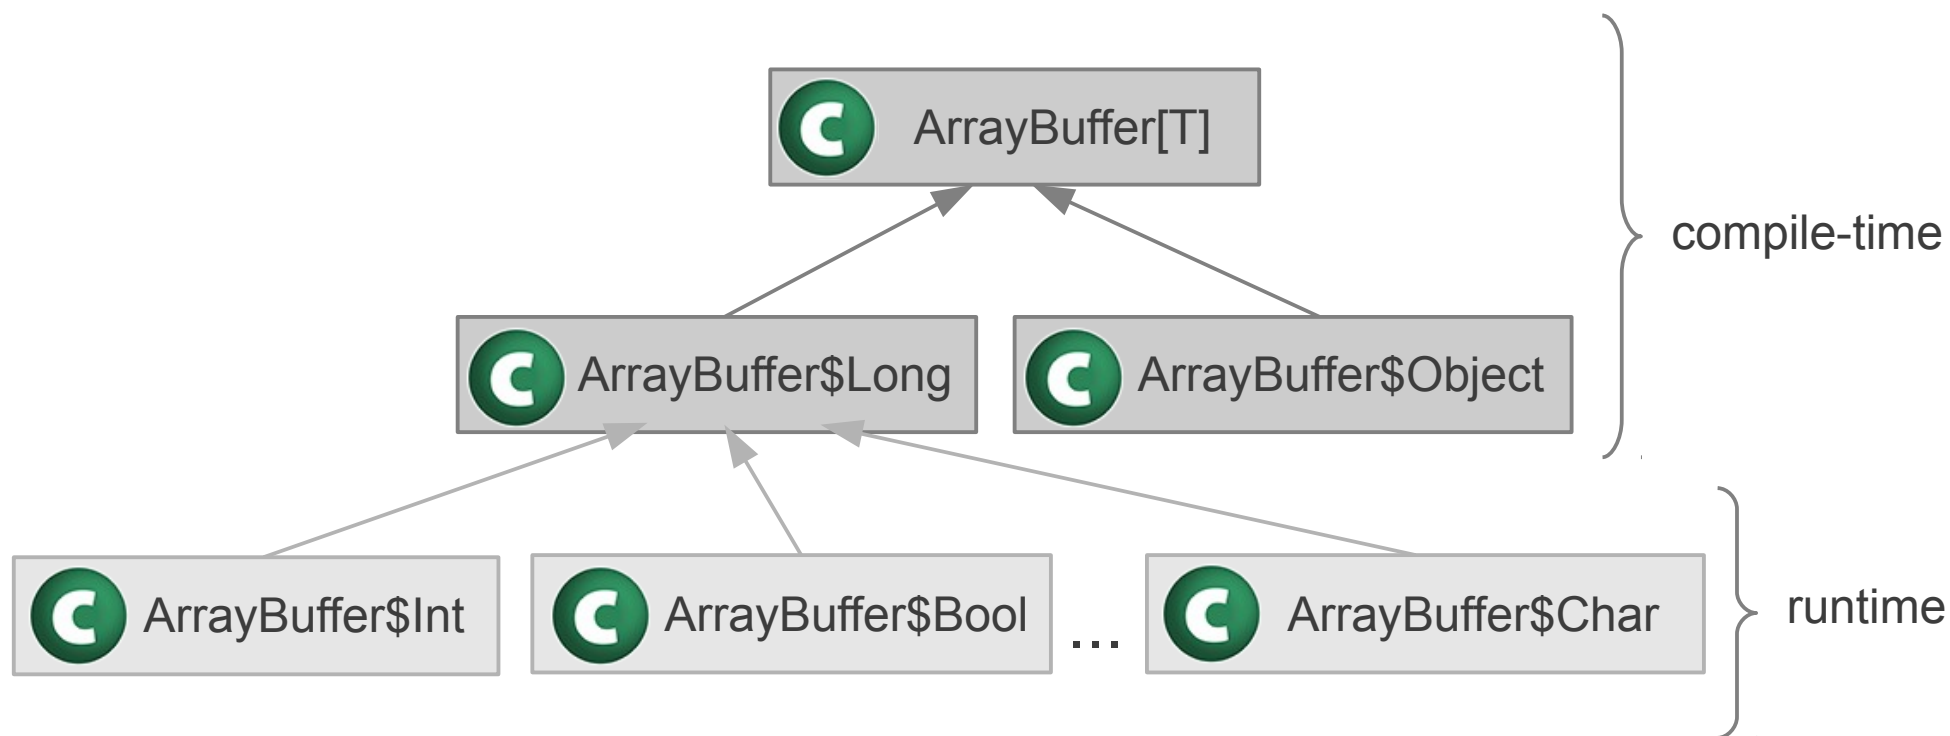


81 classes is too much - we need a lighter approach

Miniboxing?




- dot net-like runtime specialization
 - two stages: compile-time and runtime



Code is also faster than specialization

Stack Insight




- Only applied to stack values
- We'll look at arrays separately.

One **large** value type is enough

Miniboxing



- object-specialized class
- long-specialized class
- let's leave runtime aside for now



Even this is tricky. Why?

ArrayBuffer

- Uses an underlying array
 - `Array[String]`, `Array[Int]`, `Array[Float]`




Access the array through a library

Library Support

- Executing code on value types (.hashCode...)
- Interfacing with arrays

`array(i)`



```
T$type match {  
  case INT => array.asInstanceOf[Array[Int]](i).toLong  
  ...  
}
```

Surprisingly, casting and converting to long is fast

Library Support

- `ArrayBuffer.reverse()`

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

`T$type match {
 case INT => ...
 ...
}`

Can the JVM optimize such code?

Java Virtual Machine Support

- We confuse the JVM heuristics by inlining manually
- The JVM will do its best
 - For one value of T\$Type (one type)
 - hoist the switch outside the loop (loop unswitching)
 - get rid of multiple casts (common subexpression elim)
 - get rid of multiple bounds checks (check elimination)
 - **good speed**
 - For more types... there are tricks
 - **no good solution to cover all cases**

Can we lift the switch ourselves?


Dispatching

- Systems approach: T\$type
 - byte value
 - encodes the type of a parameter (INT, ...)
 - switch to do the right operation
- Object-oriented approach
 - common operations object (Dispatcher)
 - specialized instances for all types

Dispatching

- `T$Dispatcher: Dispatcher`
 - `IntDispatcher`, `LongDispatcher`, ...
- Common operators

`array(i)`



`T$Dispatcher.array_get(array, i)`

So far so good

Dispatching

- `ArrayBuffer.reverse()`

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val other = length - index - 1  
    val tmp1: T = array(index)  
    val tmp2: T = array(other)  
    array(index) = tmp2  
    array(other) = tmp1  
    index += 1  
  }  
}
```

`T$Dispatcher.array_get`

`T$Dispatcher.array_get`

`T$Dispatcher.array_update`

`T$Dispatcher.array_update`

We committed to the data representation before

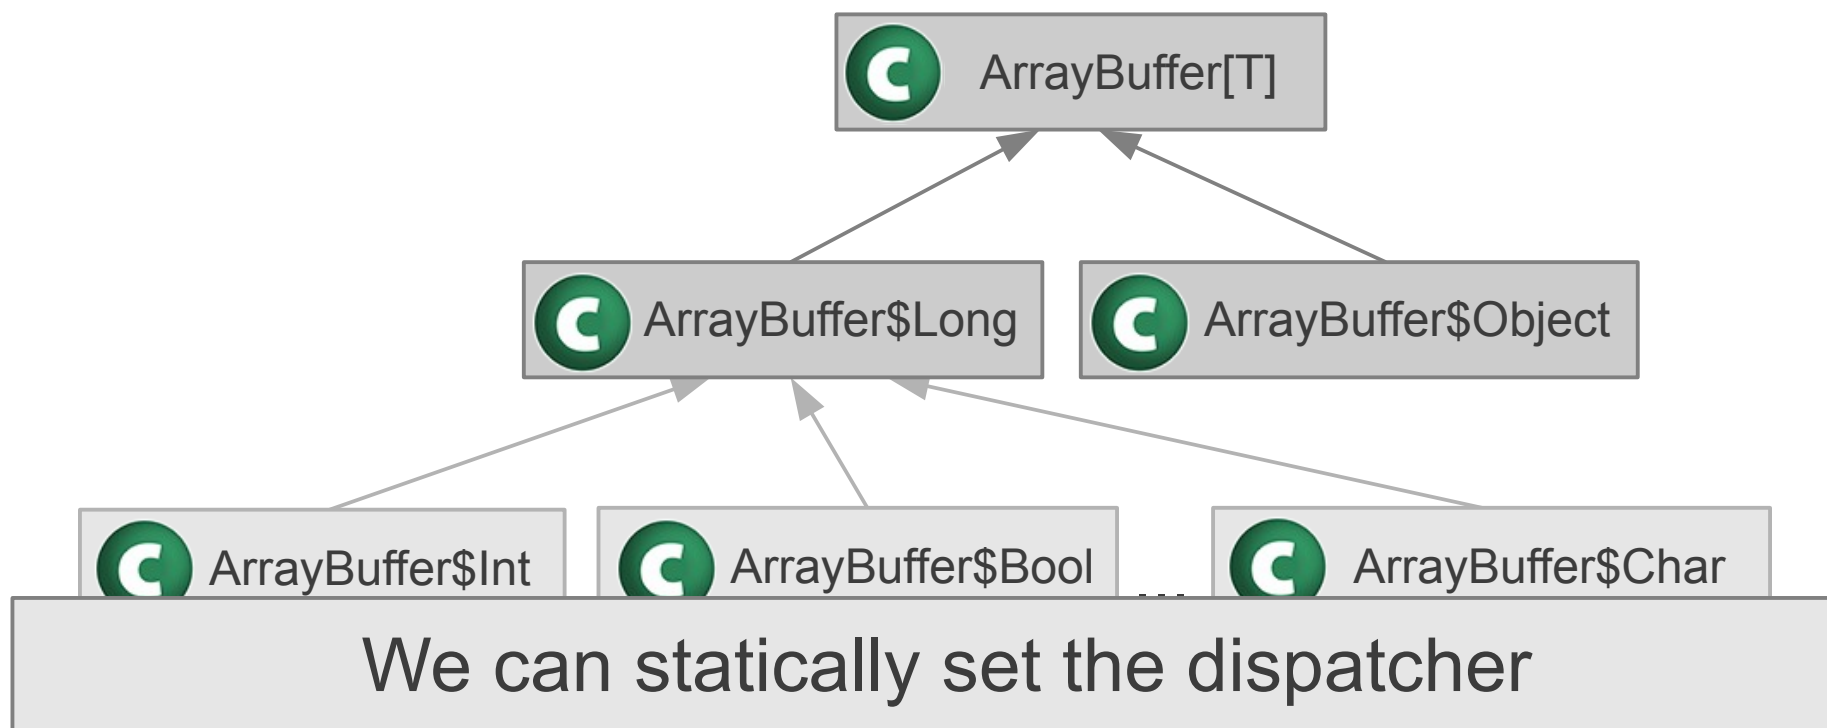
Dispatching Performance

- One or two dispatchers in a callsite
 - as fast as manually specialized code
- Three+ dispatchers
 - call to `array_get` becomes **megamorphic**
 - it is not inlined anymore
 - no optimization in the loop
 - awful results, in some cases slower than generic

Can't use such a solution. Or can we?

Runtime Specialization

- In practice, dispatcher set per INSTANCE
- Semantically - dispatcher per SPECIALIZATION
- Remember the ArrayBuffer diagram?



Runtime Specialization

- `ArrayBuffer$Int`
 - has `IntDispatcher` set statically
- Operations can be inlined
 - since `IntDispatcher` is final
 - calls never become megamorphic
 - optimizations can be done inside the loop

Class modification and loading cost?

Runtime Specialization

- One-time non-negligible cost
 - load ...\$Long bytecode
 - modify it
 - build a `Class[_]` object
 - instantiate it
- Apply the cost to the factory
 - instead of the individual class
 - turns out the cost amortizes well (10-50 uses)

Class modification and loading cost?

Numbers Preview

- `ArrayBuffer.reverse()`

ideal :	0.85883 ms +/-	0.05687
miniboxed standard mono :	0.62966 ms +/-	0.21116
miniboxed standard mega :	4.21423 ms +/-	1.15425
miniboxed standard w/cl :	0.87526 ms +/-	0.05992
miniboxed dispatch mono :	0.75230 ms +/-	0.21551
miniboxed dispatch mega :	4.90191 ms +/-	1.52366
miniboxed dispatch w/cl mono :	0.77984 ms +/-	0.14872
miniboxed dispatch w/cl mega :	0.80932 ms +/-	0.14671
specialized mono :	0.75813 ms +/-	0.16000
specialized mega :	0.61374 ms +/-	0.20144
generic :	9.76250 ms +/-	1.55297

Conclusions

- Yes, this can be done
 - There's no silver bullet
 - We just build around limitations
 - And need to make things as robust as possible
- Contributions
 - Miniboxing itself, in an open-world assumption
 - Exploration of the implementation space
 - Dispatchers
 - Runtime specialization
 - Solution that works in practice




Thank you!

github.com/miniboxing



Vlad Ureche
vlad.ureche@epfl.ch

Compatibility



- a common set of operations
 - not only the generic operations
 - we want values too

