

Miniboxing

Load-time Specialization on the JVM

OOPSLA, 29th of October 2013



Vlad Ureche
Cristian Talau
Martin Odersky

We all like generics

We all like generics

a trivial example

We all like generics

a trivial example

```
def identity[T](t: T): T = t
```

We all like generics

a trivial example

```
def identity[T](t: T): T = t
```

- will take any type and
- will return **that same type**

We all like generics

a trivial example

```
def identity[T](t: T): T = t
```

- will take any type and
- will return **that same type**

We all like generics

a trivial example

```
def identity[T](t: T): T = t
```

but under **erasure**:

```
def identity(t: Any): Any = t
```

We all like generics

a trivial example

```
def identity[T](t: T): T = t
```

but under **erasure**:

Any is the top of the
Scala type system

```
def identity(t: Any): Any = t
```


We all like generics

a trivial example

```
x = identity(3)
```

We all like generics

a trivial example

```
x = identity(3)
```

under **erasure**:

```
x = unbox(identity(box(3)))
```

**We all like generics
but under erasure**

We all like generics but under erasure

generics execute similarly to **dynamic
languages**

We all like generics but under erasure

generics execute similarly to **dynamic languages**

- generic values **lose** their **type information**

We all like generics but under erasure

generics execute similarly to **dynamic languages**

- generic values **lose** their **type information**
- primitives need **boxing**

We all like generics but under erasure

generics execute similarly to **dynamic languages**

- generic values **lose** their **type information**
- primitives need **boxing**
- performance is **affected**

We all like generics but under erasure

generics execute similarly to **dynamic languages**

- generic values **lose** their **type information**
- primitives need **boxing**
- performance is **affected**

Dynamic language VMs use **specialization** to improve performance*

We all like generics but under erasure

generics execute similarly to **dynamic languages**

- generic values **lose** their **type information**
- primitives need **boxing**
- performance is **affected**

Dynamic language VMs use **specialization** to improve performance*

*but the HotSpot JVM **doesn't**



Generics

Specialization

Miniboxing

Performance

Evaluation

WE ARE HERE



Scala has a solution

Scala has a solution

it's called specialization*

* Iulian Dragos – PhD thesis, EPFL, 2010

Scala has a solution

it's called specialization*

* Iulian Dragos – PhD thesis, EPFL, 2010

Compile-time (static) transformation

- **duplicates** the original code
- **adapts it** for each primitive type
- **rewrites** programs to use the adapted code

Scala has a solution

it's called specialization*

* Iulian Dragos – PhD thesis, EPFL, 2010

Compile-time (static) transformation

- **duplicates** the original code
- **adapts it** for each primitive type
- **rewrites** programs to use the adapted code

Adapted code doesn't need to box

Scala has a solution

it's called specialization*

*Iulian Dragos – PhD thesis, EPFL, 2010

Compile-time (static) transformation

- **duplicates** the original code
- **adapts it** for each primitive type
- **rewrites** programs to use the adapted code

Adapted code doesn't need to box

Performance is regained.

Specialization

let's revisit ``def identity``

Specialization

let's revisit `def identity`

```
def identity[T](t: T): T = t
```

Specialization

let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_V(t: Unit): Unit = t
```

```
def identity_Z(t: Boolean): Boolean = t
```

```
def identity_B(t: Byte): Byte = t
```

```
def identity_C(t: Char): Char = t
```

```
def identity_S(t: Short): Short = t
```

```
def identity_I(t: Int): Int = t
```

```
def identity_J(t: Long): Long = t
```

```
def identity_F(t: Float): Float = t
```

```
def identity_D(t: Double): Double = t
```

Specialization

let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_V(t: Unit): Unit = t
def identity_Z(t: Boolean): Boolean = t
def identity_B(t: Byte): Byte = t
def identity_C(t: Char): Char = t
def identity_S(t: Short): Short = t
def identity_I(t: Int): Int = t
def identity_L(t: Long): Long = t
def identity_F(t: Float): Float = t
def identity_D(t: Double): Double = t
```

Generates 10 times the original code

Specialization

... it gets even worse

Specialization

... it gets even worse

```
def pack[T1, T2](t1: T1, t2: T2) = ...
```

Specialization

... it gets even worse

```
def pack[T1, T2](t1: T1, t2: T2) = ...
```

```
def pack_VV(t1: Unit, t2: Unit)
def pack_VZ(t1: Unit, t2: Boolean)
def pack_VB(t1: Unit, t2: Byte)
def pack_VC(t1: Unit, t2: Char)
def pack_VS(t1: Unit, t2: Short)
def pack_VI(t1: Unit, t2: Int)
def pack_VJ(t1: Unit, t2: Long)
def pack_VF(t1: Unit, t2: Float)
def pack_VD(t1: Unit, t2: Double)
```

Specialization

... it gets even worse

```
def pack[T1, T2](t1: T1, t2: T2) = ...
```

```
def pack_VV(t1: Unit, t2: Unit)
```

```
def pack_VZ(t1: Unit, t2: Boolean)
```

```
def
```

```
def
```

10^n , where n is the number of type params

```
def pack_VS(t1: Unit, t2: Short)
```

```
def pack_VI(t1: Unit, t2: Int)
```

```
def pack_VJ(t1: Unit, t2: Long)
```

```
def pack_VF(t1: Unit, t2: Float)
```

```
def pack_VD(t1: Unit, t2: Double)
```

Specialization

... it gets even worse

```
def pack[T1, T2](t1: T1, t2: T2) = ...
```

```
def pack_VV(t1: Unit, t2: Unit)
```

```
def pack_VZ(t1: Unit, t2: Boolean)
```

```
def
```

10^n , where n is the number of type params

```
def
```

```
def
```

And this is common: Maps, Tuples, Functions

```
def
```

```
def pack_VJ(t1: Unit, t2: Long)
```

```
def pack_VF(t1: Unit, t2: Float)
```

```
def pack_VD(t1: Unit, t2: Double)
```


Specialization

... it gets even worse

```
def pack[T1, T2](t1: T1, t2: T2) = ...
```

```
def pack_VV(t1: Unit, t2: Unit)
```

```
def pack_VZ(t1: Unit, t2: Boolean)
```

def pack_VP(t1: Unit, t2: Boolean, p: Boolean) = ...
10^n, where n is the number of type params

And this is common for Lists, Tuples, Functions

```
def pack_VJ(t1: Unit, t2: Long)
```

```
def pack_VF(t1: Unit, t2: Float)
```

```
def pack_VD(t1: Unit, t2: Double)
```



Generics

Specialization

Miniboxing

WE ARE HERE

Performance

Evaluation



Miniboxing




Miniboxing

reduces the variants



Miniboxing

reduces the variants



by using something like a **tagged union**




Miniboxing

reduces the variants




by using something like a **tagged union**



Miniboxing

reduces the variants



by using something like a **tagged union**



Miniboxing

reduces the variants



by using something like a **tagged union**



Miniboxing

reduces the variants



by using something like a **tagged union**



Miniboxing

reduces the variants




by using something like a **tagged union**

| | TAG | DATA (VALUE) |
|---------|------|--------------|
| false = | BOOL | 0x0 |
| true = | BOOL | 0x1 |

Miniboxing

reduces the variants



by using something like a **tagged union**

| | TAG | DATA (VALUE) |
|---------|------|--------------|
| false = | BOOL | 0x0 |
| true = | BOOL | 0x1 |
| 42 = | INT | 0x2A |

Miniboxing

reduces the variants




by using something like a **tagged union**



Miniboxing

reduces the variants




by using something like a **tagged union**



and using the **static type information**

- tags are attached to **code**, not to values

Miniboxing



let's revisit `def identity`

```
def identity[T](t: T): T = t
```

Miniboxing



let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_M(T_tag: Byte, t: Long): Long
```

Miniboxing




let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_M(T_tag: Byte, t: Long): Long
```




Miniboxing



let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_M(T_tag: Byte, t: Long): Long
```




Miniboxing



let's revisit `def identity`


```
def identity[T](t: T): T = t
```

```
def identity_M(T_tag: Byte, t: Long): Long
```



T_tag corresponds to the **type parameter**, instead of the values being passed around.


Miniboxing



let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_M(T_tag: Byte, t: Long): Long
```



`T_tag` corresponds to the **type parameter**, instead of the value bound.

Tag hoisting

Miniboxing




let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_M(T_tag: Byte, t: Long): Long
```

Two variants per type parameter (reference + minibox)

Miniboxing



let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_M(T_tag: Byte, t: Long): Long
```

Two variants per type parameter (reference + minibox)

`def pack` will have 4 variants

Miniboxing



let's revisit `def identity`

```
def identity[T](t: T): T = t
```

```
def identity_M(T_tag: Byte, t: Long): Long
```

Two variants per type parameter (reference + minibox)

`def pack` will have 4 variants

Tag hoisting is instrumental in
obtaining good performance



Generics

Specialization

Miniboxing

Performance

WE ARE HERE

Evaluation




Performance

needs one more ingredient

$$\boxed{\text{Miniboxing}} = \boxed{\text{Tagged union}}$$

Performance

needs one more ingredient



Performance

needs one more ingredient



Performance

needs one more ingredient




Why do we need a **secret ingredient**?

Switching on tags kills performance



Switching on tags kills performance



```
def toString(T_tag: Byte,  
             value: Long): String =  
  T_tag match {  
    case UNIT => ...  
    case BOOL => ...  
    ...  
  }
```


Switching on tags kills performance



```
def toString(T_tag: Byte,  
             value: Long): String =  
  T_tag match {  
    case UNIT => ...  
    case BOOL => ...  
    ...  
  }
```

Even more so for consecutive switches

Switching on tags kills performance



```
T_tag match {  
  case X => op1  
}  
T_tag match {  
  case X => op2  
}
```

Switching on tags kills performance



```
T_tag match {  
  case X => op1  
}  
T_tag match {  
  case X => op2  
}
```



Switching on tags kills performance



```
T_tag match {  
  case X => op1  
}  
T_tag match {  
  case X => op2  
}
```

Redundant
switch

Switching on tags kills performance



```
T_tag match {  
  case X => op1; op2  
}  
T_tag match {  
  case X =>  
}
```

← Redundant
switch

Switching on tags kills performance





```
T_tag match {  
  case X => op1; op2  
}
```



```
T_tag match {  
  case X =>  
}
```

Redundant
switch

Switching on tags kills performance




Switching on tags kills performance



This is critical for **array operations**


Switching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

Switching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

T_tag match {
 case INT => ...
 ...
}

Switching




ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

T_tag match {
 case INT => ...
 ..
}

T_tag match {
 case INT => ...
 ...
}

Switching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

```
T_tag match {  
  case INT => ...  
  ..  
}
```

```
T_tag match {  
  case INT => ...
```

```
T_tag match {  
  case INT => ...
```

```
T_tag match {  
  case INT => ...  
  ...  
}
```

Switching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

T_tag match {
 case INT => ...
 ..
}


T_tag match {
 case INT => ...

T_tag match {
 case INT => ...

T_tag match {
 case INT => ...
 ...
}

Fuse the operations together?

Switching




ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

```
T_tag match {  
  case INT =>  
    val tmp1 = ...  
    val tmp2 = ...  
    array(.) = ...  
    array(.) = ...  
    ...  
}
```

Switching




ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

```
T_tag match {  
  case INT =>  
    val tmp1 = ...  
    val tmp2 = ...  
    array(.) = ...  
    array(.) = ...  
    ...  
}
```

Hoist the switch out of the loop?

Switching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

```
T_tag match {  
  case INT =>  
    var index = 0  
    while (...) {  
      ...  
      index += 1  
    }  
}
```

Switching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(opposite)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

```
T_tag match {  
  case INT =>  
    var index = 0  
    while (...) {  
      ...  
      index += 1  
    }  
}
```

Is that enough? Method may be **called from a loop**

Performance

needs one more ingredient



Performance

needs one more ingredient



Can't be switching

Performance

needs one more ingredient



Can't be switching


Must be something else

Dispatching

- Dispatch object
 - Encodes array interactions



Dispatching



- Dispatch object
 - Encodes array interactions

```
class Dispatcher[T] {  
  def array_get(...): Long  
  def array_set(...): Unit  
}
```

Dispatching



- Dispatch object
 - Encodes array interactions

```
class Dispatcher[T] {  
  def array_get(...): Long  
  def array_set(...): Unit  
}
```

```
def identity_M(T_dispatcher: Dispatcher[T],  
              t: Long): Long
```

Dispatching




- Dispatch object
 - Encodes array interactions

```
class Dispatcher[T] {  
  def array_get(...): Long  
  def array_set(...): Unit  
}
```

```
def identity_M(T_dispatcher: Dispatcher[T],  
              t: Long): Long
```

instead of tag


Dispatching



- Dispatch object
 - Encodes array interactions

```
object IntDispatcher extends Dispatcher[Int] {  
  def array_get(...): Long = ...  
  def array_set(...): Unit = ...  
}
```

Dispatching




- Dispatch object
 - Encodes array interactions

```
object IntDispatcher extends Dispatcher[Int] {  
  def array_get(...): Long = ...  
  def array_set(...): Unit = ...  
}
```

```
object LongDispatcher ...  
object CharDispatcher ...
```

Dispatching




- Dispatch object
 - Encodes array interactions

```
object IntDispatcher extends Dispatcher[Int] {  
  def array_get(...): Long = ...  
  def array_set(...): Unit = ...  
}
```

```
object LongDispatcher ...  
object CharDispatcher ...
```

Passing a dispatcher = hoisted already

Dispatching



`ArrayBuffer.reverse()`

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(other)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

Dispatching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(other)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

T_dispatcher.array_get

Dispatching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(other)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

T_dispatcher.array_get

T_dispatcher.array_get

Dispatching



ArrayBuffer.reverse()

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length - index - 1  
    val tmp1: T = array(index)  
    val tmp2: T = array(other)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

T_dispatcher.array_get

T_dispatcher.array_get

T_dispatcher.array_set

T_dispatcher.array_set

Dispatching



`ArrayBuffer.reverse()`

```
def reverse(): Unit {  
  var index = 0  
  while (index * 2 < length) {  
    val opposite = length-index-1  
    val tmp1: T = array(index)  
    val tmp2: T = array(other)  
    array(index) = tmp2  
    array(opposite) = tmp1  
    index += 1  
  }  
}
```

`T_dispatcher.array_get`

`T_dispatcher.array_get`

`T_dispatcher.array_set`

`T_dispatcher.array_set`

With inlining, we get good performance

Dispatching

```
ArrayBuffer.reverse()
```

```
T_dispatcher.array_get
```



Dispatching




`ArrayBuffer.reverse()`

`T_dispatcher.array_get`

`IntDispatcher`

Monomorphic, okay

Dispatching



`ArrayBuffer.reverse()`

`T_dispatcher.array_get`

`IntDispatcher`

`LongDispatcher`

Monomorphic, okay

Polymorphic, okay

Dispatching



`ArrayBuffer.reverse()`

`T_dispatcher.array_get`

IntDispatcher

LongDispatcher

DoubleDispatcher

Monomorphic, okay

Polymorphic, okay

Megamorphic* → no more inlining

*** for the HotSpot JVM**

Dispatching



`ArrayBuffer.reverse()`

`T_dispatcher.array_get`

IntDispatcher

LongDispatcher

DoubleDispatcher

Monomorphic, okay

Polymorphic, okay

Megamorphic* → no more inlining

*** for the HotSpot JVM**

No more inlining → bad performance

Performance

needs one more ingredient



Performance

needs one more ingredient

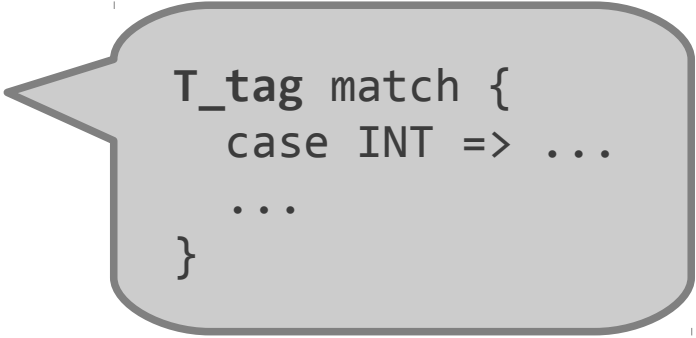


Object oriented dispatch isn't that

The secret ingredient

The secret ingredient

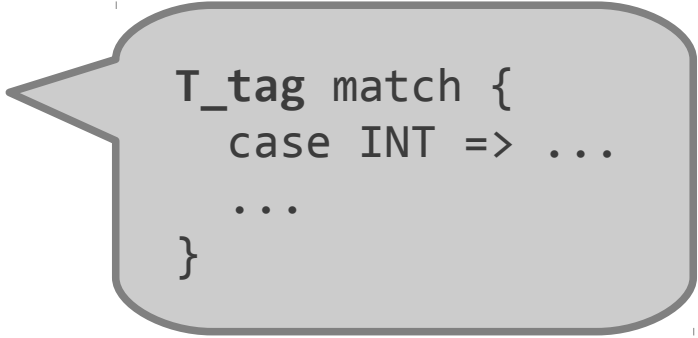
- **Switch-based** dispatching



```
T_tag match {  
  case INT => ...  
  ...  
}
```

The secret ingredient

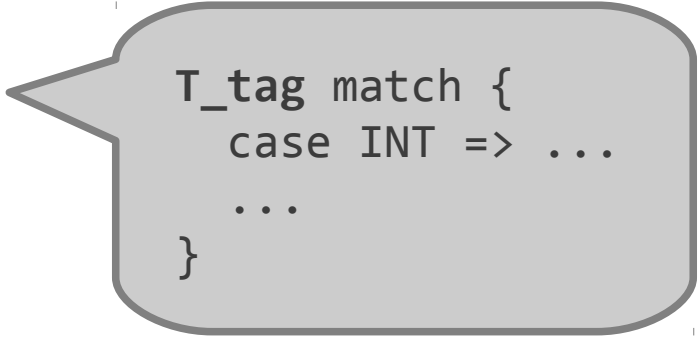
- **Switch-based** dispatching
- When instantiating the class
 - **T_tag** is known



```
T_tag match {  
  case INT => ...  
  ...  
}
```

The secret ingredient

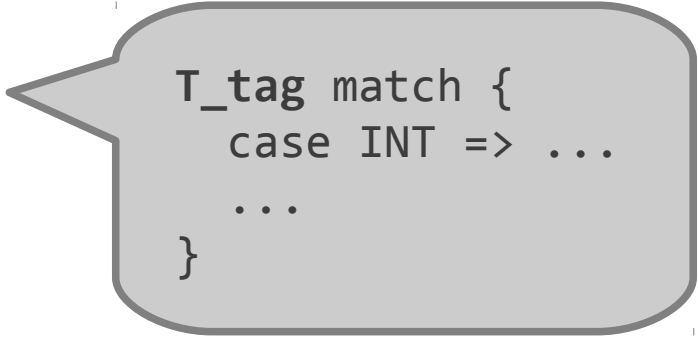
- **Switch-based** dispatching
- When instantiating the class
 - **T_tag** is **known**
 - **T_tag** is a **constant**



```
T_tag match {  
  case INT => ...  
  ...  
}
```


The secret ingredient

- **Switch-based** dispatching
- When instantiating the class
 - **T_tag** is **known**
 - **T_tag** is a **constant**

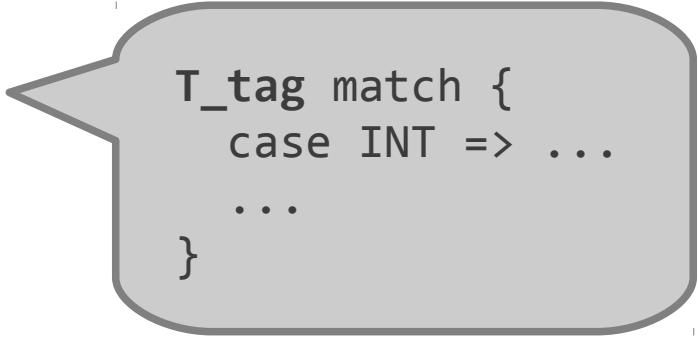


```
T_tag match {  
  case INT => ...  
  ...  
}
```

Encode T_tag in the class name?

The secret ingredient

- **Switch-based** dispatching
- When instantiating the class
 - **T_tag** is **known**
 - **T_tag** is a **constant**



```
T_tag match {  
  case INT => ...  
  ...  
}
```

Encode T_tag in the class name?

Statically? Code explosion!

Load-time specialization

- Load-time transformation

```
T_tag match {  
  case INT => ...  
  case CHAR => ...  
  case UNIT => ...  
  ...  
}
```

Load-time specialization

- Load-time transformation
 - set `T_tag` statically

```
INT match {  
  case INT => ...  
  case CHAR => ...  
  case UNIT => ...  
  ...  
}
```

Load-time specialization

- Load-time transformation
 - set `T_tag` statically
 - perform **constant folding**

```
INT match {  
  case INT => ...  
  case CHAR => ...  
  case UNIT => ...  
  ...  
}
```

Load-time specialization

- Load-time transformation
 - set `T_tag` statically
 - perform **constant folding**
 - perform **dead code elimination**



Load-time specialization

- Load-time transformation
 - set `T_tag` statically
 - perform **constant folding**
 - perform **dead code elimination**



Only the useful code

Load-time specialization

- Load-time transformation
 - set `T_tag` statically
 - perform **constant folding**
 - perform **dead code elimination**



Only the useful code

No dispatching

Load-time specialization

- Load-time transformation
 - set `T_tag` statically
 - perform **constant folding**
 - perform **dead code elimination**

...

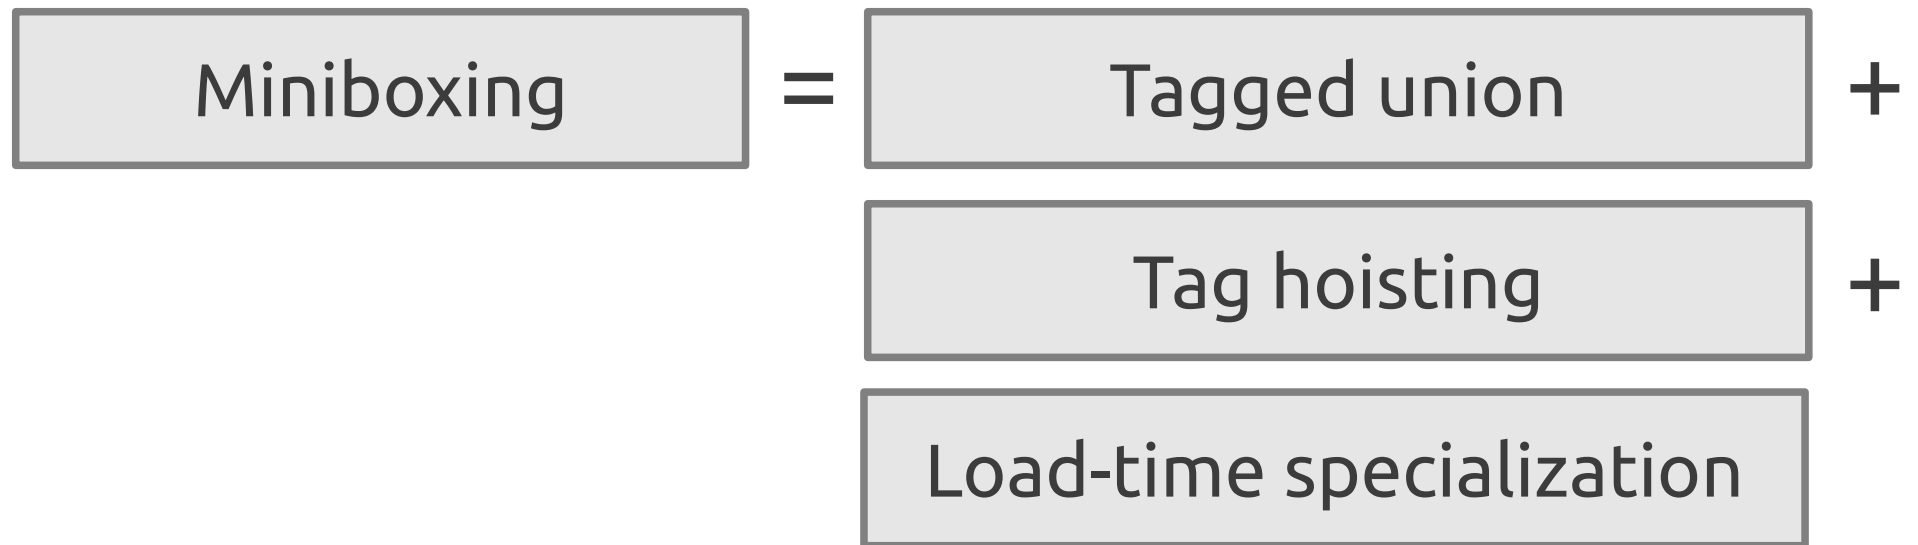
Only the useful code

No dispatching

Is this the secret ingredient? **Yes!**


Performance

needs one more ingredient



Performance

needs one more ingredient



Attaching tags to code enables
load-time specialization



Generics

Specialization


Miniboxing

Performance


Evaluation




Evaluation - Performance




Evaluation - Performance




Evaluation - Performance



Evaluation - Performance




Evaluation - Performance



Similar results on other benchmarks


Evaluation - Code size

Spire – numeric abstractions library (12KLOC)



Evaluation - Code size

Spire – numeric abstractions library (12KLOC)



2.8x bytecode reduction (4.7x for Vector in std. lib)

Contributions



Miniboxing

Contributions

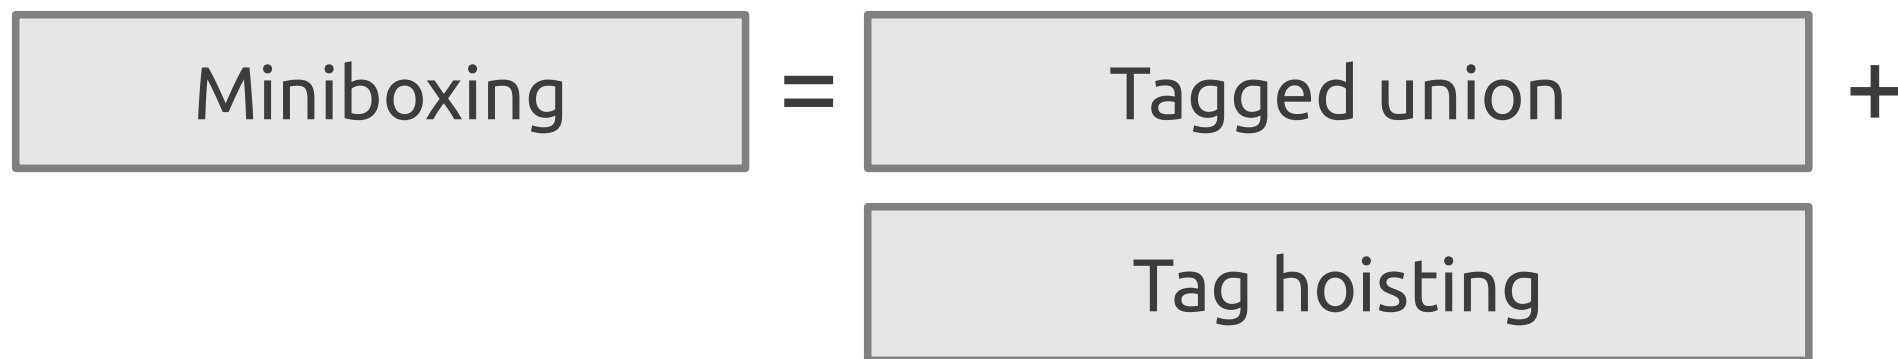


Miniboxing

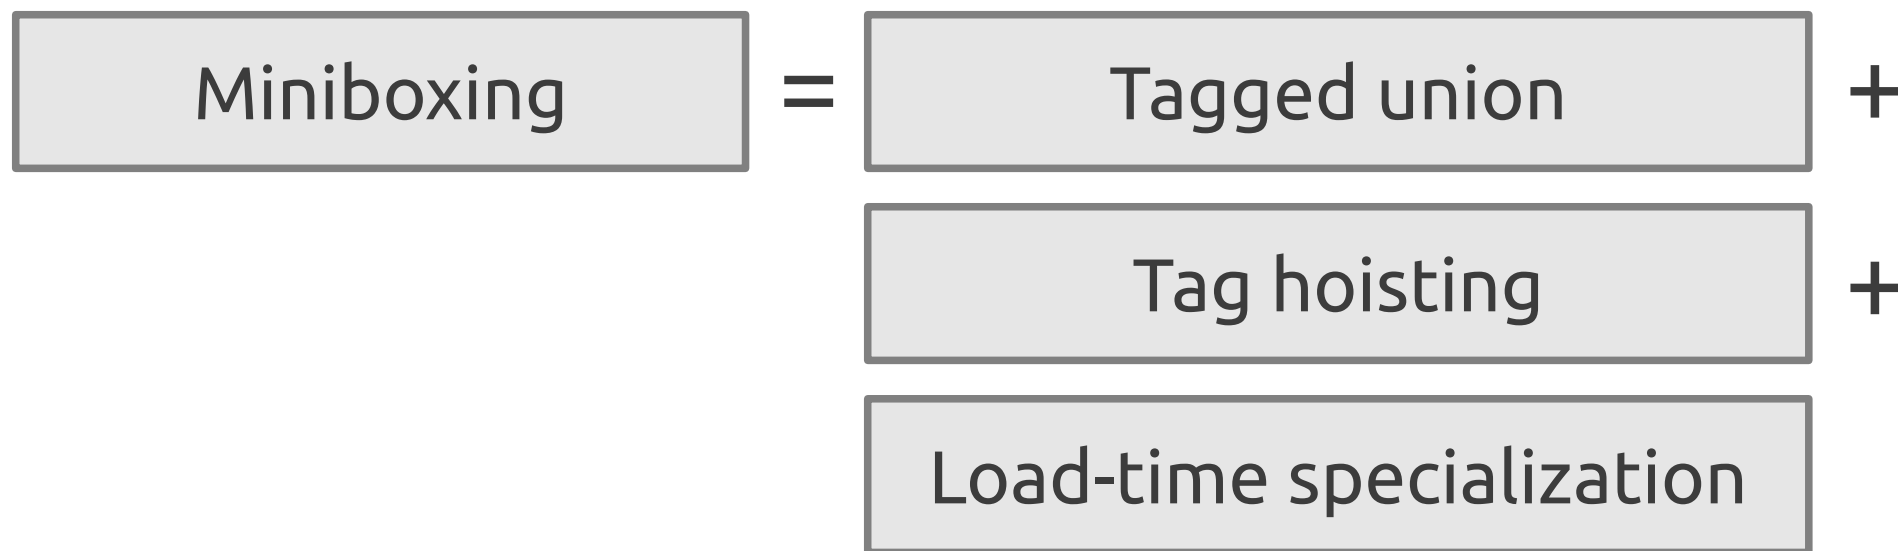

=

Tagged union


Contributions



Contributions



Conclusions



- improves performance
- reduces bytecode size

visit scala-miniboxing.org!